

▼ 1 - Fully Connected Network from scratch

▼ 1.1 Loading the FashionMNIST dataset

▼ Download the dataset

The source comes from Fashion-MNIST <https://github.com/zalandoresearch/fashion-mnist>

Note: If dataset is already downloaded, it is not downloaded again.

```
import torchvision

transform = torchvision.transforms.Compose(
    [torchvision.transforms.ToTensor(), torchvision.transforms.Normalize(
        (0.5,0.5,0.5), (0.5,0.5,0.5))])
trainset = torchvision.datasets.FashionMNIST(root='./data', train=True,
                                              download=True, transform=transform)
testset = torchvision.datasets.FashionMNIST(root='./data', train=False,
                                              download=True, transform=transform)
```

▼ Pre-processing the dataset

Get the training dataset.

```
X_train = trainset.data
print(X_train.shape)
```

```
↳ torch.Size([60000, 28, 28])
```

Flatten (60000,28,28) size of X_train into (784,60000) size of matrix.

Also, normalize each pixel value to between 0 and 1.

```
import numpy as np
X_train = X_train.reshape(X_train.shape[0], -1).T / 255.0
print(X_train.shape)
```

```
↳
```

Get the testing dataset

```
X_test = testset.data  
print(X_test.shape)
```



Flatten (10000,28,28) size of X_test into (784,10000) size of matrix.

Also, normalize each pixel value to between 0 and 1.

```
X_test = X_test.reshape(X_test.shape[0], -1).T / 255.0  
print(X_test.shape)
```



Get the label from training dataset.

```
y_train = trainset.targets  
print(y_train)
```



Transfer training label from tensor to numpy.

```
y_train = np.reshape(np.asarray(y_train), (60000,))  
print(y_train.shape)
```



Get the label from testing dataset.

And, transfer it into numpy.

```
y_test = testset.targets  
y_test = np.reshape(np.asarray(y_test), (10000,))  
print(y_test.shape)
```



This is the label class:

```
label = ('T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
```

```
'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle Boot']
```

Let's take a look of the first two picture we get in the training dataset:

```
fp1 = X_train[:,0]
fp2 = X_train[:,1]

import matplotlib.pyplot as plt
f, ax = plt.subplots(1,2)

ax[0].imshow(fp1.reshape((28,28)), cmap="gray") #first image
ax[1].imshow(fp2.reshape((28,28)), cmap="gray") #second image
plt.show()
```



The label of the training dataset shown above, the first picture is index 9 and second picture is index 0.

As we can see in the label class, index 9 is 'Ankle Boot', index 0 is 'T-shirt/top'.

▼ Shuffle the dataset

Shuffle the training and validation dataset, with label, by index.

```
def split(X, y, val_ratio=0.1):
    val_number = int(val_ratio * X.shape[1])
    random_indice = np.random.permutation(X.shape[1])
    return X[:, random_indice[val_number:]], y[random_indice[val_number:]], \
           X[:, random_indice[:val_number]], \
           y[random_indice[:val_number]]

X_train, y_train, X_val, y_val = split(X_train, y_train)
```

▼ One Hot Encoding Labels

Make one dimension labels from (dims,) to (10, dims) for the better prediction of the Neural Network.

```
def one_hot_encoding(label):
    shape = (label.max() + 1, label.size)
    rows = np.arange(label.size)
    one_hot = np.zeros(shape)
    one_hot[label, rows] = 1
    return one_hot

y_train = one_hot_encoding(y_train)
y_val = one_hot_encoding(y_val)
y_test = one_hot_encoding(y_test)

print(y_test.shape)
```



▼ 1.2 Implementing the Network

```
class NeuralNetwork:
    def __init__(self, layer_dims, seed_value=99):
        """
        Arguments:
        layer_dims -- A list contains the dimensions of each layer in CNN
        seed_value -- set seed to generate random numpy

        Attributes generated:
        parameters -- a dict contains parameters "W1", "b1", ..., "WL",
                       "bL" of each corresponding layer
                       W1 -- weight matrix of shape (layer_dims[1], layer_dims[1-1])
                       b1 -- bias vector of shape (layer_dims[1], 1)
        nn_architecture -- the dictionary contains the description of the
                           network architecture
                           'input_dim' -- input dimension from last network layer
                           'output_dim' -- output dimension after applying the
                                           activation
                           'activation' -- ReLU or Softmax
        X_val -- validation data in training
        y_val -- validation label in training
        cost_history -- stores cross entropy loss for every training iteration
        acc_history -- stores the training accuracy for every iteration
        """

        np.random.seed(seed_value)
        self.parameters = {}
        self.nn_architecture = []
```

```

self.X_val = None
self.y_val = None
self.cost_history = []
self.acc_history = []

input_dim = layer_dims[0]
for idx, output_dim in enumerate(layer_dims[1:]):
    self.nn_architecture.append({'input_dim': input_dim, 'output_dim': \
                                output_dim, 'activation': 'relu'})
    input_dim = output_dim
self.nn_architecture[len(self.nn_architecture) - 1]['activation'] = \
    'softmax'

for idx, layer in enumerate(self.nn_architecture):
    layer_idx = idx + 1
    layer_input_size = layer["input_dim"]
    layer_output_size = layer["output_dim"]

    self.parameters['W' + str(layer_idx)] = np.random.randn(
        layer_output_size, layer_input_size) * 0.1
    self.parameters['b' + str(layer_idx)] = np.random.randn(
        layer_output_size, 1) * 0.1

def relu(self, Z):
    """
    Implement the ReLU function

    Arguments:
    Z -- output of linear portion of the output layer (last layer)

    Returns:
    Rectified linear units
    """
    return np.maximum(0, Z)

def softmax(self, Z):
    """
    Implement the Softmax function

    Arguments:
    Z -- output of linear portion of the output layer (last layer)

    Returns:
    p -- softmax probability
    """
    p = np.exp(Z - np.max(Z, axis=0, keepdims=True))
    p /= np.sum(p, axis=0, keepdims=True)
    return p

def affineForward(self, A, W, b, activation="relu"):
    """
    Implement the linear portion of CNN's forward propagation

```

Implement the linear portion of CNN's forward propagation.

Arguments:

A -- activation from previous layer
W -- weights matrix
b -- bias vector

Returns:

The neural layer after applied activation function

Z -- the input of the activation function

"""

Z = np.dot(W, A) + b

if activation == "relu":

 activation_func = self.relu

elif activation == "softmax":

 activation_func = self.softmax

else:

 raise Exception('Non-supported activation function')

return activation_func(Z), Z

def forwardPropagation(self, X):

"""

Implement the forward propagation

Arguments:

X -- input from input layer as the starting point of forward propagation

Returns:

A_curr -- the output of the activation function after each layer

cache -- a list stores parameters (cache) for each layer during
forward propagation

"""

A_curr = X

cache = {}

for idx, layer in enumerate(self.nn_architecture):

 layer_idx = idx + 1

 A_prev = A_curr

 activation_function_curr = layer["activation"]

 W_curr = self.parameters["W" + str(layer_idx)]

 b_curr = self.parameters["b" + str(layer_idx)]

 A_curr, Z_curr = self.affineForward(A_prev, W_curr, b_curr,
 activation_function_curr)

 cache["A" + str(idx)] = A_prev

 cache["Z" + str(layer_idx)] = Z_curr

return A_curr, cache

```

def costFunction(self, AL, y, epsilon=1e-12):
    """
    Implement the cross entropy loss

    Arguments:
    AL -- output of linear portion of the output layer (last layer)
    y -- the labels of data
    epsilon -- a very small value, in case the AL is zero when log

    Returns:
    cost -- the cross entropy loss value of each iteration of forward
            propagation
    """
    cost = -np.mean(np.multiply(y, np.log(AL + epsilon)))
    return cost

def derivative_cost(self, AL, y):
    """
    Implement the first step of back propagation: the derivative of cost
    function over AL

    Arguments:
    AL -- output of linear portion of the output layer (last layer)
    y -- the labels of data

    Returns:
    dAL -- the derivative of cost function over AL
    """
    dAL = AL - y
    return dAL

def affineBackward(self, dA_curr, W_curr, Z_curr, A_prev, activation="relu"):
    """
    Implement the linear portion of backward propagation of one layer(1)

    Arguments:
    dA_curr -- Gradient of the cost over activation output from current layer (layer l-1)
    W_curr -- W for the current layer(layer1)
    Z_curr -- Z for the current layer(layer1)
    A_prev -- A for the previous layer(layer1)

    Returns:
    dA_prev -- Gradient of the cost over activation output from previous layer (layer l-1)
    dW_curr -- Gradient of the cost over W for the current layer(layer1)
    db_curr -- Gradient of the cost over b for the current layer(layer1)
    """
    if activation == "relu":
        dZ_curr = self.derivative_relu(dA_curr, Z_curr)
    elif activation == "softmax":
        dZ_curr = dA_curr
    else:
        raise Exception('Non-supported activation function')

```

```

raise Exception('non supported activation function')

m = A_prev.shape[1]
dW_curr = np.dot(dZ_curr, A_prev.T) / m
db_curr = np.sum(dZ_curr, axis=1, keepdims=True) / m
dA_prev = np.dot(W_curr.T, dZ_curr)

return dA_prev, dW_curr, db_curr

def derivative_relu(self, dA, Z):
    """
    Implement the derivative calculation of relu activation during backpropagation

    Arguments:
    dA -- the gradient of cost function over relu activation output A

    Returns:
    dZ -- the derivative of cost function over Z
    """
    dZ = np.array(dA, copy=True)
    dZ[Z <= 0] = 0
    return dZ

def backPropagation(self, dAL, cache):
    """
    Implement backpropagation for each layer

    Arguments:
    dAL -- the gradient of cost function over the activation output of last
           layer: starting point for backpropagation
    cache -- a list stores tuple of A,W,b,Z for each layer

    Returns:
    grads -- a dict stores dWl and dbl for each layer l
    """
    grads = {}

    # initiation of backward propagation
    dA_prev = dAL

    for layer_idx_prev, layer in reversed(list(enumerate(
                                                self.nn_architecture))):
        # we number network layers from 1
        layer_idx_curr = layer_idx_prev + 1
        # extraction of the activation function for the current layer
        activ_function_curr = layer["activation"]
        # update dA
        dA_curr = dA_prev
        # get variables
        A_prev = cache["A" + str(layer_idx_prev)]
        Z_curr = cache["Z" + str(layer_idx_curr)]
        W_curr = self.parameters["W" + str(layer_idx_curr)]

```



```

    # get the gradient of the cost
    dA_prev, dW_curr, db_curr = self.affineBackward(
        dA_curr, W_curr, Z_curr, A_prev, activ_function_curr)
    # update gradient values
    grads["dW" + str(layer_idx_curr)] = dW_curr
    grads["db" + str(layer_idx_curr)] = db_curr

return grads

def updateParameters(self, grads, alpha):
    """
    Use gradient descent to implement updated parameters

    Arguments:
    grads -- a dict stores all parameters gradients for each layer
    alpha -- learning rate
    """
    for layer_idx, layer in enumerate(self.nn_architecture, 1):
        self.parameters["W" + str(layer_idx)] -= alpha * grads["dW" + \
                                                                str(layer_idx)]
        self.parameters["b" + str(layer_idx)] -= alpha * grads["db" + \
                                                                str(layer_idx)]

def train(self, X, y, iters, alpha, batch_size, verbose=False, print_every=100):
    """
    It takes advantage of every function in this class to implement training
    and validation using CNN.

    Arguments:
    X -- input data
    y -- labels of data
    iters -- number of iterations to run
    alpha -- learning rate
    batch_size -- number of samples to assign to minibatch
    verbose -- decide whether to print or not print loss and accuracy
    print_every -- number of iterations to print

    Return:
    Number of iterations, train loss, train_acc, and valid_acc in every
    customized iterations
    """
    for i in range(0, iters):
        X_batch, y_batch = self.get_batch(X, y, batch_size)
        AL, cache = self.forwardPropagation(X_batch)
        loss = self.costFunction(AL, y_batch)
        dAL = self.derivative_cost(AL, y_batch)
        grads = self.backPropagation(dAL, cache)
        self.updateParameters(grads, alpha)
        train_acc = self.score(self.predict(X), y)
        if verbose:
            if i == 0 or (i+1) % print_every == 0:
                val_acc = self.score(self.predict(self.X_val), self.y_val)

```

```

        print('iter={:4}, loss={:.6f}, train_acc={:.6f}, validation_acc={:.6f}'.format(
            self.cost_history.append(loss)
            self.acc_history.append(train_acc)

def predict(self, X):
    """
    It predicts the label given input x
    Argument:
    X -- input of data
    Return:
    y_pred -- predicted label
    """
    AL, _ = self.forwardPropagation(X)
    y_pred = np.argmax(AL, axis=0)
    return y_pred

def score(self, y_pred, y_true):
    """
    It calculates the percentage of correct predicted labels over true labels

    Argument:
    y_pred-- predicted labels
    y_true -- true labels

    Return: percentage of correct predicted labels
    """
    target = np.argmax(y_true, axis=0)
    correct = np.mean(y_pred == target)
    return correct

def load_validation_set(self, X_val, y_val):
    """
    Load validation set to CNN
    """
    self.X_val = X_val
    self.y_val = y_val

def get_batch(self, X, y, batch_size):
    """
    Load mini-batch to CNN
    """
    batch_index = np.random.randint(X.shape[1], size=batch_size)
    X_batch = X[:, batch_index]
    y_batch = y[:, batch_index]
    return X_batch, y_batch

```

▼ 1.3 Predicting

Define variables for training the Fully Connected Neural Network:

```
batch_size = 1000
seed_value = 2
learning_rate = 0.1
iteration = 120
print_every = 10
layer_dims = [X_train.shape[0], 1024, 256, 128, 10]
```

Train the F-CNN:

```
CNN = NeuralNetwork(layer_dims, seed_value)
CNN.load_validation_set(X_val, y_val)
CNN.train(X_train, y_train, iteration, learning_rate, batch_size, True, print_every)
```

```
↳ iter= 1, loss=0.429308, train_acc=0.100111, validation_acc=0.099167
   iter= 10, loss=0.166481, train_acc=0.461574, validation_acc=0.455000
   iter= 20, loss=0.101982, train_acc=0.656241, validation_acc=0.646500
   iter= 30, loss=0.096968, train_acc=0.680870, validation_acc=0.676167
   iter= 40, loss=0.071267, train_acc=0.709759, validation_acc=0.699333
   iter= 50, loss=0.070132, train_acc=0.763889, validation_acc=0.761167
   iter= 60, loss=0.067600, train_acc=0.742463, validation_acc=0.736667
   iter= 70, loss=0.061559, train_acc=0.771593, validation_acc=0.767667
   iter= 80, loss=0.059510, train_acc=0.788759, validation_acc=0.784000
   iter= 90, loss=0.051820, train_acc=0.806685, validation_acc=0.804333
   iter= 100, loss=0.068209, train_acc=0.765444, validation_acc=0.759833
   iter= 110, loss=0.053147, train_acc=0.820815, validation_acc=0.819500
   iter= 120, loss=0.063519, train_acc=0.789130, validation_acc=0.785333
```


Test the accuracy based on the testing dataset:

```
y_pred = CNN.predict(X_test)
test_acc = CNN.score(y_pred, y_test)
print('Accuracy predicting test dataset = {:.2%}'.format(test_acc))
```

```
↳ Accuracy predicting test dataset = 77.97%
```


To visualize the performance of the Neural Network, a plot below shows the cross entropy loss during the iterations:

```
import matplotlib.pyplot as plt
plt.plot(CNN.cost_history)
plt.xlabel('Iterations')
plt.ylabel('Cross entropy loss')
plt.show()
```



Below shows the prediction accuracy of the training data:

```
plt.plot(CNN.acc_history)
plt.xlabel('Iterations')
plt.ylabel('Accuracy')
plt.show()
```



▼ 2 - CNN using PyTorch

▼ Libraries used:

```
import torch
import torch.nn as nn
from torch.autograd import Variable
import torch.nn.functional as F
```

```
import torchvision
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms

import matplotlib.pyplot as plt
```

▼ 2.1 Loading the FashionMNIST dataset

▼ Download the dataset

```
trainSet = torchvision.datasets.FashionMNIST(root='./data', train=True, download=True,
                                             transform=transforms.ToTensor())
testSet = torchvision.datasets.FashionMNIST(root='./data', train=False, download=True,
                                             transform=transforms.ToTensor())
```

Make sure you turn on the GPU in Runtime setting, for a better performance.

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

▼ Get the train/test loader using DataLoader

With batch size --> 100, and Shuffle is TRUE.

```
trainLoader = torch.utils.data.DataLoader(trainSet, batch_size=100, shuffle=True)
testLoader = torch.utils.data.DataLoader(testSet, batch_size=100, shuffle=True)
```

▼ Creating a method to name the class for the label.

example: 9 --> Ankle Boot

```
def outputLabel(label):
    outputMapping = {
        0: "T-shirt/Top",
        1: "Trouser",
        2: "Pullover",
        3: "Dress",
        4: "Coat",
        5: "Sandal",
        6: "Shirt",
        7: "Sneaker",
```

```

        8: "Bag",
        9: "Ankle Boot"
    }
    input = (label.item() if type(label) == torch.Tensor else label)
    return outputMapping[input]


```

▼ To show a data using matplotlib :

```

image, label = next(iter(trainSet))
plt.imshow(image.squeeze(), cmap="gray")
print(label)

```



▼ 2.2 Implementing the Network

▼ Building the CNN class:

Model Class Name: FashionCNN

Layers: 2 sequential layers consist of:

```

* Convolution layer with kernal -> 3*3, padding = 1 (1st layer) & padding = 0 (2nd layer).
* Stride of 1 in both the layer
* Activation function: Relu
* All the functionaltiy is given in forward method that defines the forward pass of CNN.
* Output:
    1. 1st Conv later:    input: 28 * 28 * 3    and  Output : 28 * 28 * 32.
    2. Max pooling layer: input: 28 * 28 * 32   and Output:  14 * 14 * 32.
    3. 2nd Conv layer:    input : 14 * 14 * 32  and output: 12 * 12 * 64
    4. 2nd Max Pooling layer : 12 * 12 * 64, output: 6 * 6 * 64.

```

At last fully connected layer has 10 output features for 10 types of clothes.

```
class FashionCNN(nn.Module):

    def __init__(self):
        super(FashionCNN, self).__init__()

        self.layer1 = nn.Sequential(
            nn.Conv2d(in_channels=1, out_channels=32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2)
        )

        self.layer2 = nn.Sequential(
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3),
            nn.BatchNorm2d(64),
            nn.ReLU(),
            nn.MaxPool2d(2)
        )

        self.fc1 = nn.Linear(in_features=64*6*6, out_features=600)
        self.drop = nn.Dropout2d(0.25)
        self.fc2 = nn.Linear(in_features=600, out_features=120)
        self.fc3 = nn.Linear(in_features=120, out_features=10)

    def forward(self, x):
        out = self.layer1(x)
        out = self.layer2(out)
        out = out.view(out.size(0), -1)
        out = self.fc1(out)
        out = self.drop(out)
        out = self.fc2(out)
        out = self.fc3(out)

        return out
```

▼ Method 1 - Chage another CNN architecture:


This architecture is defined in Lecture 4.

With kernel sizes (5 x 5), Flattened input in fc_3 and softmax at the output.

n1, n2 and n3 in this case, are 32, 64 and 128.

The CNN architecture is shown in the picture below:

Convolutional Neural Network



(Note: if you choose to use the first FashionCNN, then don't run the following code)

```
class FashionCNN(nn.Module):
    def __init__(self):
        super(FashionCNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 5, 1)
        self.conv2 = nn.Conv2d(32, 64, 5, 1)
        self.dropout1 = nn.Dropout2d(0.5)
        self.fc3 = nn.Linear(4*4*64, 128)
        self.fc4 = nn.Linear(128, 10)

        # x represents our data

    def forward(self, x):
        # Pass data through conv1
        x = self.conv1(x)
        # Use the rectified-linear activation function over x
        x = F.relu(x)
        # Run max pooling over x
        x = F.max_pool2d(x, 2)
        # Pass data through conv2
        x = self.conv2(x)
        x = F.relu(x)
        # Run max pooling over x
        x = F.max_pool2d(x, 2)
        # Flatten x with start_dim=1
        x = torch.flatten(x, 1)
```



```

x = torch.flatten(x, 1)
# Pass data through fc3
x = self.fc3(x)
x = F.relu(x)
# With dropout
x = self.dropout1(x)
# Pass data through fc4
x = self.fc4(x)

# Apply softmax to output
output = F.log_softmax(x, dim=1)
return output

```

▼ Define CNN by using different architecture to construct model

Take a look of the chosen model:

```

model = FashionCNN()
model.to(device)
print(model)

```

```

↳ FashionCNN(
  (layer1): Sequential(
    (0): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU()
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (layer2): Sequential(
    (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU()
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (fc1): Linear(in_features=2304, out_features=600, bias=True)
  (drop): Dropout2d(p=0.25, inplace=False)
  (fc2): Linear(in_features=600, out_features=120, bias=True)
  (fc3): Linear(in_features=120, out_features=10, bias=True)
)

```

▼ Method 2 - Change different optimizers

Firstly, define the learnin rate:

```

learningRate = 0.01

```

Secondly, choose the optimizer in torch. First one we used, is ADAM:

```
optimizer = torch.optim.Adam(model.parameters(), lr=learningRate)
```

The second optimizer we used is SGD:

(Note, you only need to run either the code above or the following one.)

```
optimizer = torch.optim.SGD(model.parameters(), lr=learningRate, momentum=0.9)
```

▼ 2.3 Predicting

Let us start training the network with chosen architecture and optimizer, then test it on the testing dataset

```
error = nn.CrossEntropyLoss()
numEpochs = 15
count = 0
# Lists for visualization of loss and accuracy
lossList = []
iterationList = []
accuracyList = []

# Lists for knowing classwise accuracy
predictionsList = []
labelsList = []

for epoch in range(numEpochs):
    for images, labels in trainLoader:
        # Transferring the images and labels to GPU(if available)
        images, labels = images.to(device), labels.to(device)

        train = Variable(images.view(100, 1, 28, 28))
        labels = Variable(labels)

        # Here is --> Forward pass
        outputs = model(train)
        loss = error(outputs, labels)

        # Initializing a gradient as 0 so there is no mixing of gradient among the batches
        optimizer.zero_grad()

        # Propagating the error backward
        loss.backward()

        # Optimizing the parameters (Adam Algorithm)
        optimizer.step()
```

```
optimizer.step()

count += 1

# Let us start testing the model:

if not (count % 50):
    total = 0
    correct = 0

    for images, labels in testLoader:
        images, labels = images.to(device), labels.to(device)
        labelsList.append(labels)

        test = Variable(images.view(100, 1, 28, 28))

        outputs = model(test)

        predictions = torch.max(outputs, 1)[1].to(device)
        predictionsList.append(predictions)
        correct += (predictions == labels).sum()

        total += len(labels)
    accuracy = correct * 1.0 / total
    lossList.append(loss.data)
    iterationList.append(count)
    accuracyList.append(accuracy)


if not (count % 500):
    print("Iteration: {}, Loss: {:.3f}, Accuracy: {:.2%}".format(count, loss.data, ac
```

```
↳ Iteration: 500, Loss: 0.265, Accuracy: 86.36%
Iteration: 1000, Loss: 0.150, Accuracy: 88.40%
Iteration: 1500, Loss: 0.326, Accuracy: 89.72%
Iteration: 2000, Loss: 0.320, Accuracy: 90.50%
Iteration: 2500, Loss: 0.191, Accuracy: 90.32%
Iteration: 3000, Loss: 0.261, Accuracy: 90.05%
Iteration: 3500, Loss: 0.337, Accuracy: 90.77%
Iteration: 4000, Loss: 0.233, Accuracy: 90.64%
Iteration: 4500, Loss: 0.222, Accuracy: 90.64%
Iteration: 5000, Loss: 0.110, Accuracy: 90.47%
Iteration: 5500, Loss: 0.135, Accuracy: 91.72%
Iteration: 6000, Loss: 0.142, Accuracy: 90.32%
Iteration: 6500, Loss: 0.182, Accuracy: 90.73%
Iteration: 7000, Loss: 0.185, Accuracy: 91.04%
Iteration: 7500, Loss: 0.140, Accuracy: 91.18%
Iteration: 8000, Loss: 0.127, Accuracy: 90.62%
Iteration: 8500, Loss: 0.084, Accuracy: 91.31%
Iteration: 9000, Loss: 0.085, Accuracy: 91.31%
```

Plot the loss in iterations:


```
plt.plot(iterationList, lossList)
```

```
plt.plot(iterationList, lossList)
plt.xlabel("Iterations")
plt.ylabel("Loss")
plt.show()
```



Plot the accuracy in iterations:

```
plt.plot(iterationList, accuracyList)
plt.xlabel("Iterations")
plt.ylabel("Accuracy")
plt.show()
```



▼ Test

```
model.eval()
correct = 0
for images, labels in testLoader:
    with torch.no_grad(): # so that computation graph history is not stored
        images, labels = images.to(device), labels.to(device) # send tensors to GPU
```

```

        outputs = model(images)
        predictions = outputs.data.max(1)[1]
        correct += predictions.eq(labels.data).sum()

print('Test set accuracy: {:.2f}%'.format(100.0 * correct / len(testLoader.dataset)))

```

☞ Test set accuracy: 91.46%

Auxillary: Getting the accuracy with respect to each classification in the fashionMINST Dataset

```

classCorrect = [0. for _ in range(10)]
totalCorrect = [0. for _ in range(10)]

with torch.no_grad():
    for images, labels in testLoader:
        images, labels = images.to(device), labels.to(device)
        test = Variable(images)
        outputs = model(test)
        predicted = torch.max(outputs, 1)[1]
        c = (predicted == labels).squeeze()

        for i in range(100):
            label = labels[i]
            classCorrect[label] += c[i].item()
            totalCorrect[label] += 1

for i in range(10):
    print("Accuracy of {}: {:.2f}%".format(outputLabel(i), classCorrect[i] * 100 / totalCorrect[i]))

```

☞ Accuracy of T-shirt/Top: 82.40%
 Accuracy of Trouser: 98.50%
 Accuracy of Pullover: 84.60%
 Accuracy of Dress: 93.90%
 Accuracy of Coat: 87.10%
 Accuracy of Sandal: 98.00%
 Accuracy of Shirt: 77.60%
 Accuracy of Sneaker: 97.80%
 Accuracy of Bag: 97.60%
 Accuracy of Ankle Boot: 97.10%

Among 10 categories, Bag has the highest prediction accuracy, and Shirt has the lowest.

▼ Conclusion according to Method 1

With same other conditions, the first Architecture we used, is more instable. The loss shown in the iteration has big waves.

Instead, the second Architecture defined according to the lecture, has a little bit more smooth loss curve and more stable.

▼ Conclusion according to Method 2

For both Architecture, ADAM predicts the accuracy faster. Although, SGD is more stable in loss calculation.

Furthurmore, we found using Architecture 1 or the Architecture 2 with SGD can reach their best performance overall.

▼ 3 - Answer these questions

▼ Question 1

Convolution makes the neural network to model translation of object in image automatically, could CNN model the different orientations of object in images? How to train a CNN to classify objects with different orientations?

▼ Answer 1

Yes, the CNN model can be used for the different orientations of the object in images.

Firstly, we will look in brute force:

We can give the rotated images while training the model. However, we don't think it is an optimal solution. In fact, we believe this is a weak solution.

While there are different methods we can use, We will be explaining the identification of the different orientations of images using an available model named RotNet.

What is RotNet?

It is a convolutional neural network for predicting the rotation angle of an image that corrects its orientation.

To explain this model in detail:

We have two consecutive convolutional layers (with kernel size & number of the different kernels). The next layer is the max-pooling layer (which also takes the kernel size). The dropout layer sets a fraction of its input to zero. The flatten layer simply converts the 3-D input to 1-D.

Finally, we have a dropout layer & the fully connected layer.

In general, there is an activation function & the final layer uses Softmax activation.

Also, we use an ADAM optimizer used to perform the weights update.

▼ Question 2

Random dropout is used to regularize fully connect layers and avoid overfitting of neural networks. If the dropout rate is 0.2, describing how the dropout connection performed in training and testing (backward and forward)?

▼ Answer 2

When the dropout set to 0.2, that means in the forward propagation, 20 percent of the input data are zeroed. When backward propagation, only 80% neurons in that layer are usable, since 20% are killed.